

# **Inter-Task Communication Mechanisms**

# Global shared variables

In most RT applications, tasks exchange data through global shared variables.

## *Advantages*

- High efficiency
- Low run-time overhead
- Schedulability analysis is available

## *Disadvantages*

- Data must be accessed in mutual exclusion:
  - non-preemptive regions;
  - semaphores, priority inheritance, priority ceiling.
- Not good for modular design:
  - local details are exposed to other tasks;
  - a change in a task affect other tasks.

# Message passing paradigm

Another approach is to exchange data through a **message passing paradigm**:

- Every task operates on a private memory space;
- Data are exchanged by messages through a channel.



**Message:** set of data having a predefined format.

**Channel:** logical link by which two tasks can communicate.

# Communication Ports

- Many operating systems provides the channel abstraction through the **port** construct.
- A task can use a port to exchange messages by means of two primitives:

***send***      sends a message to a port

***receive***    receives a message from a port

Some operating systems allow the user to define different types of ports, with peculiar semantics.

# Port characteristics


Ports may differ for:

- number of tasks allowed to send messages;
- number of tasks allowed to receive messages;
- **policy** used to insert and extract messages.
- behavior to manage **exceptions** (sending to a full port or receiving from an empty port).

- 
- Before being used, a port has to be **created**, and then **destroyed** when it is not needed any more.
  - Port **attributes** need to be defined at creation time

# Sending a message

A message sent to a port is inserted to an internal buffer, whose size must be defined at creation time:




If a message is sent when the port buffer is full, an exception policy has to be selected. Typically:

- the sender is blocked until the receiver reads a message (synchronous behavior);
- the new message is lost;
- an error message is returned by the send.

# Receiving a message

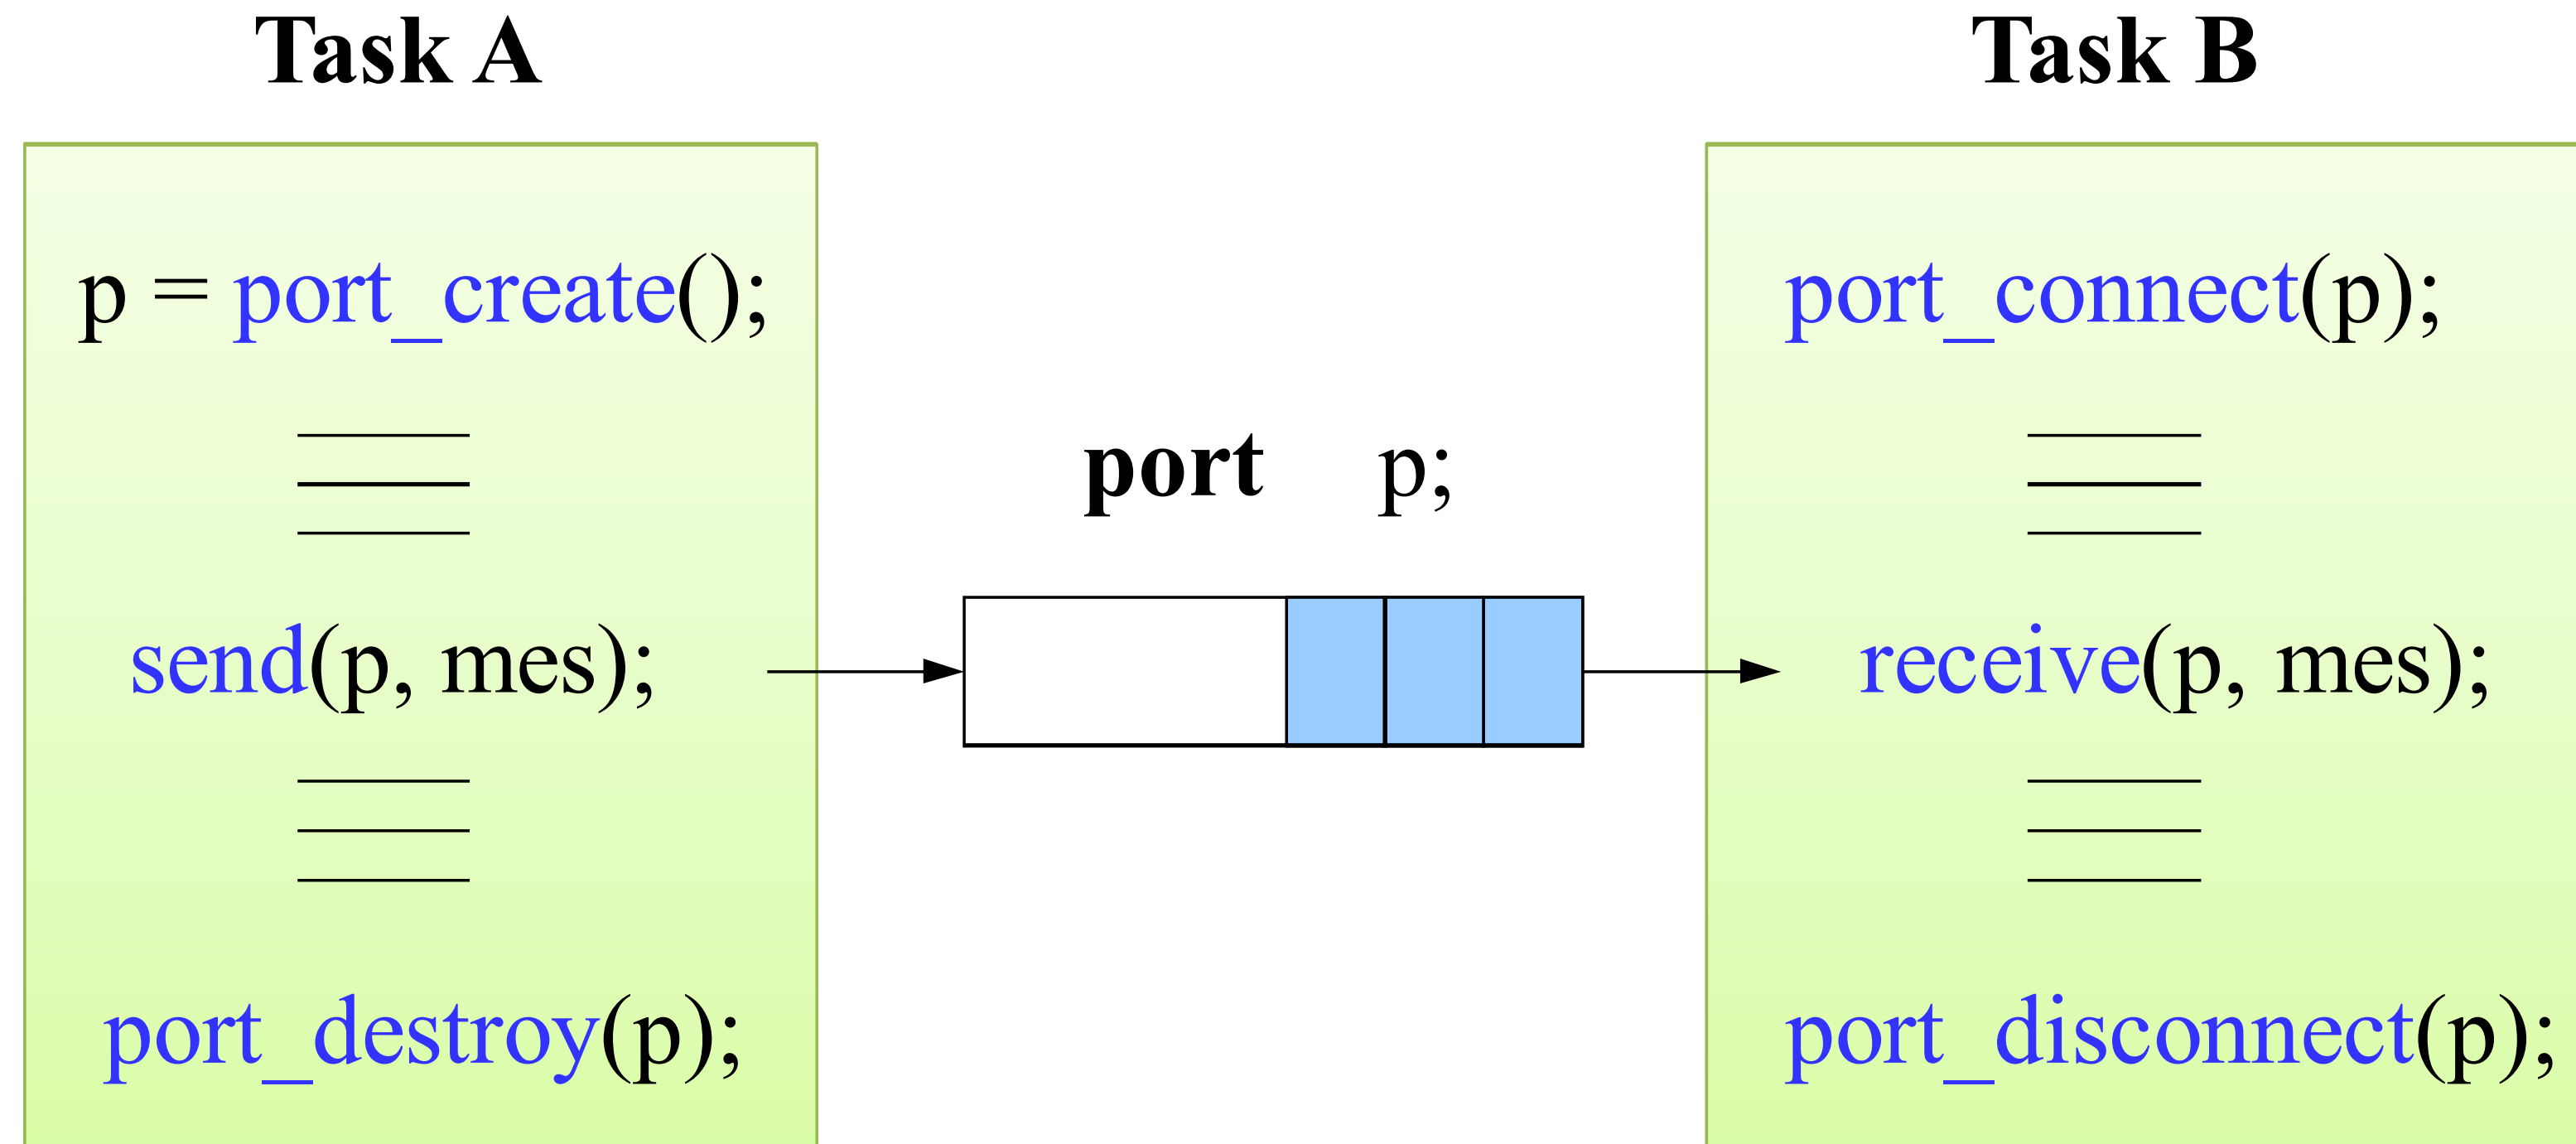
When receiving from a port, the message at the head of the buffer is extracted (consumed):



When receiving from a port with an empty buffer, an exception policy has to be selected. Typically:

- the receiver is blocked until a new message is sent (synchronous behavior);
- an error message is returned by the receive.


# Using a port



**NOTE:** Task A is the owner and must start first.




# Periodic task communication



## Problem ( $T_1 < T_2$ )

- If ( $T_1 < T_2$ ),  $\tau_1$  puts more messages than what  $\tau_2$  can read.
- When the buffer becomes full,  $\tau_1$  must proceed at the same rate as  $\tau_2$ .

# Periodic task communication




## Problem ( $T_1 > T_2$ )

- If ( $T_1 > T_2$ ),  $\tau_2$  reads more messages than what  $\tau_1$  can produce.
- When the buffer becomes empty,  $\tau_2$  must proceed at the same rate as  $\tau_1$ .

# Periodic task communication

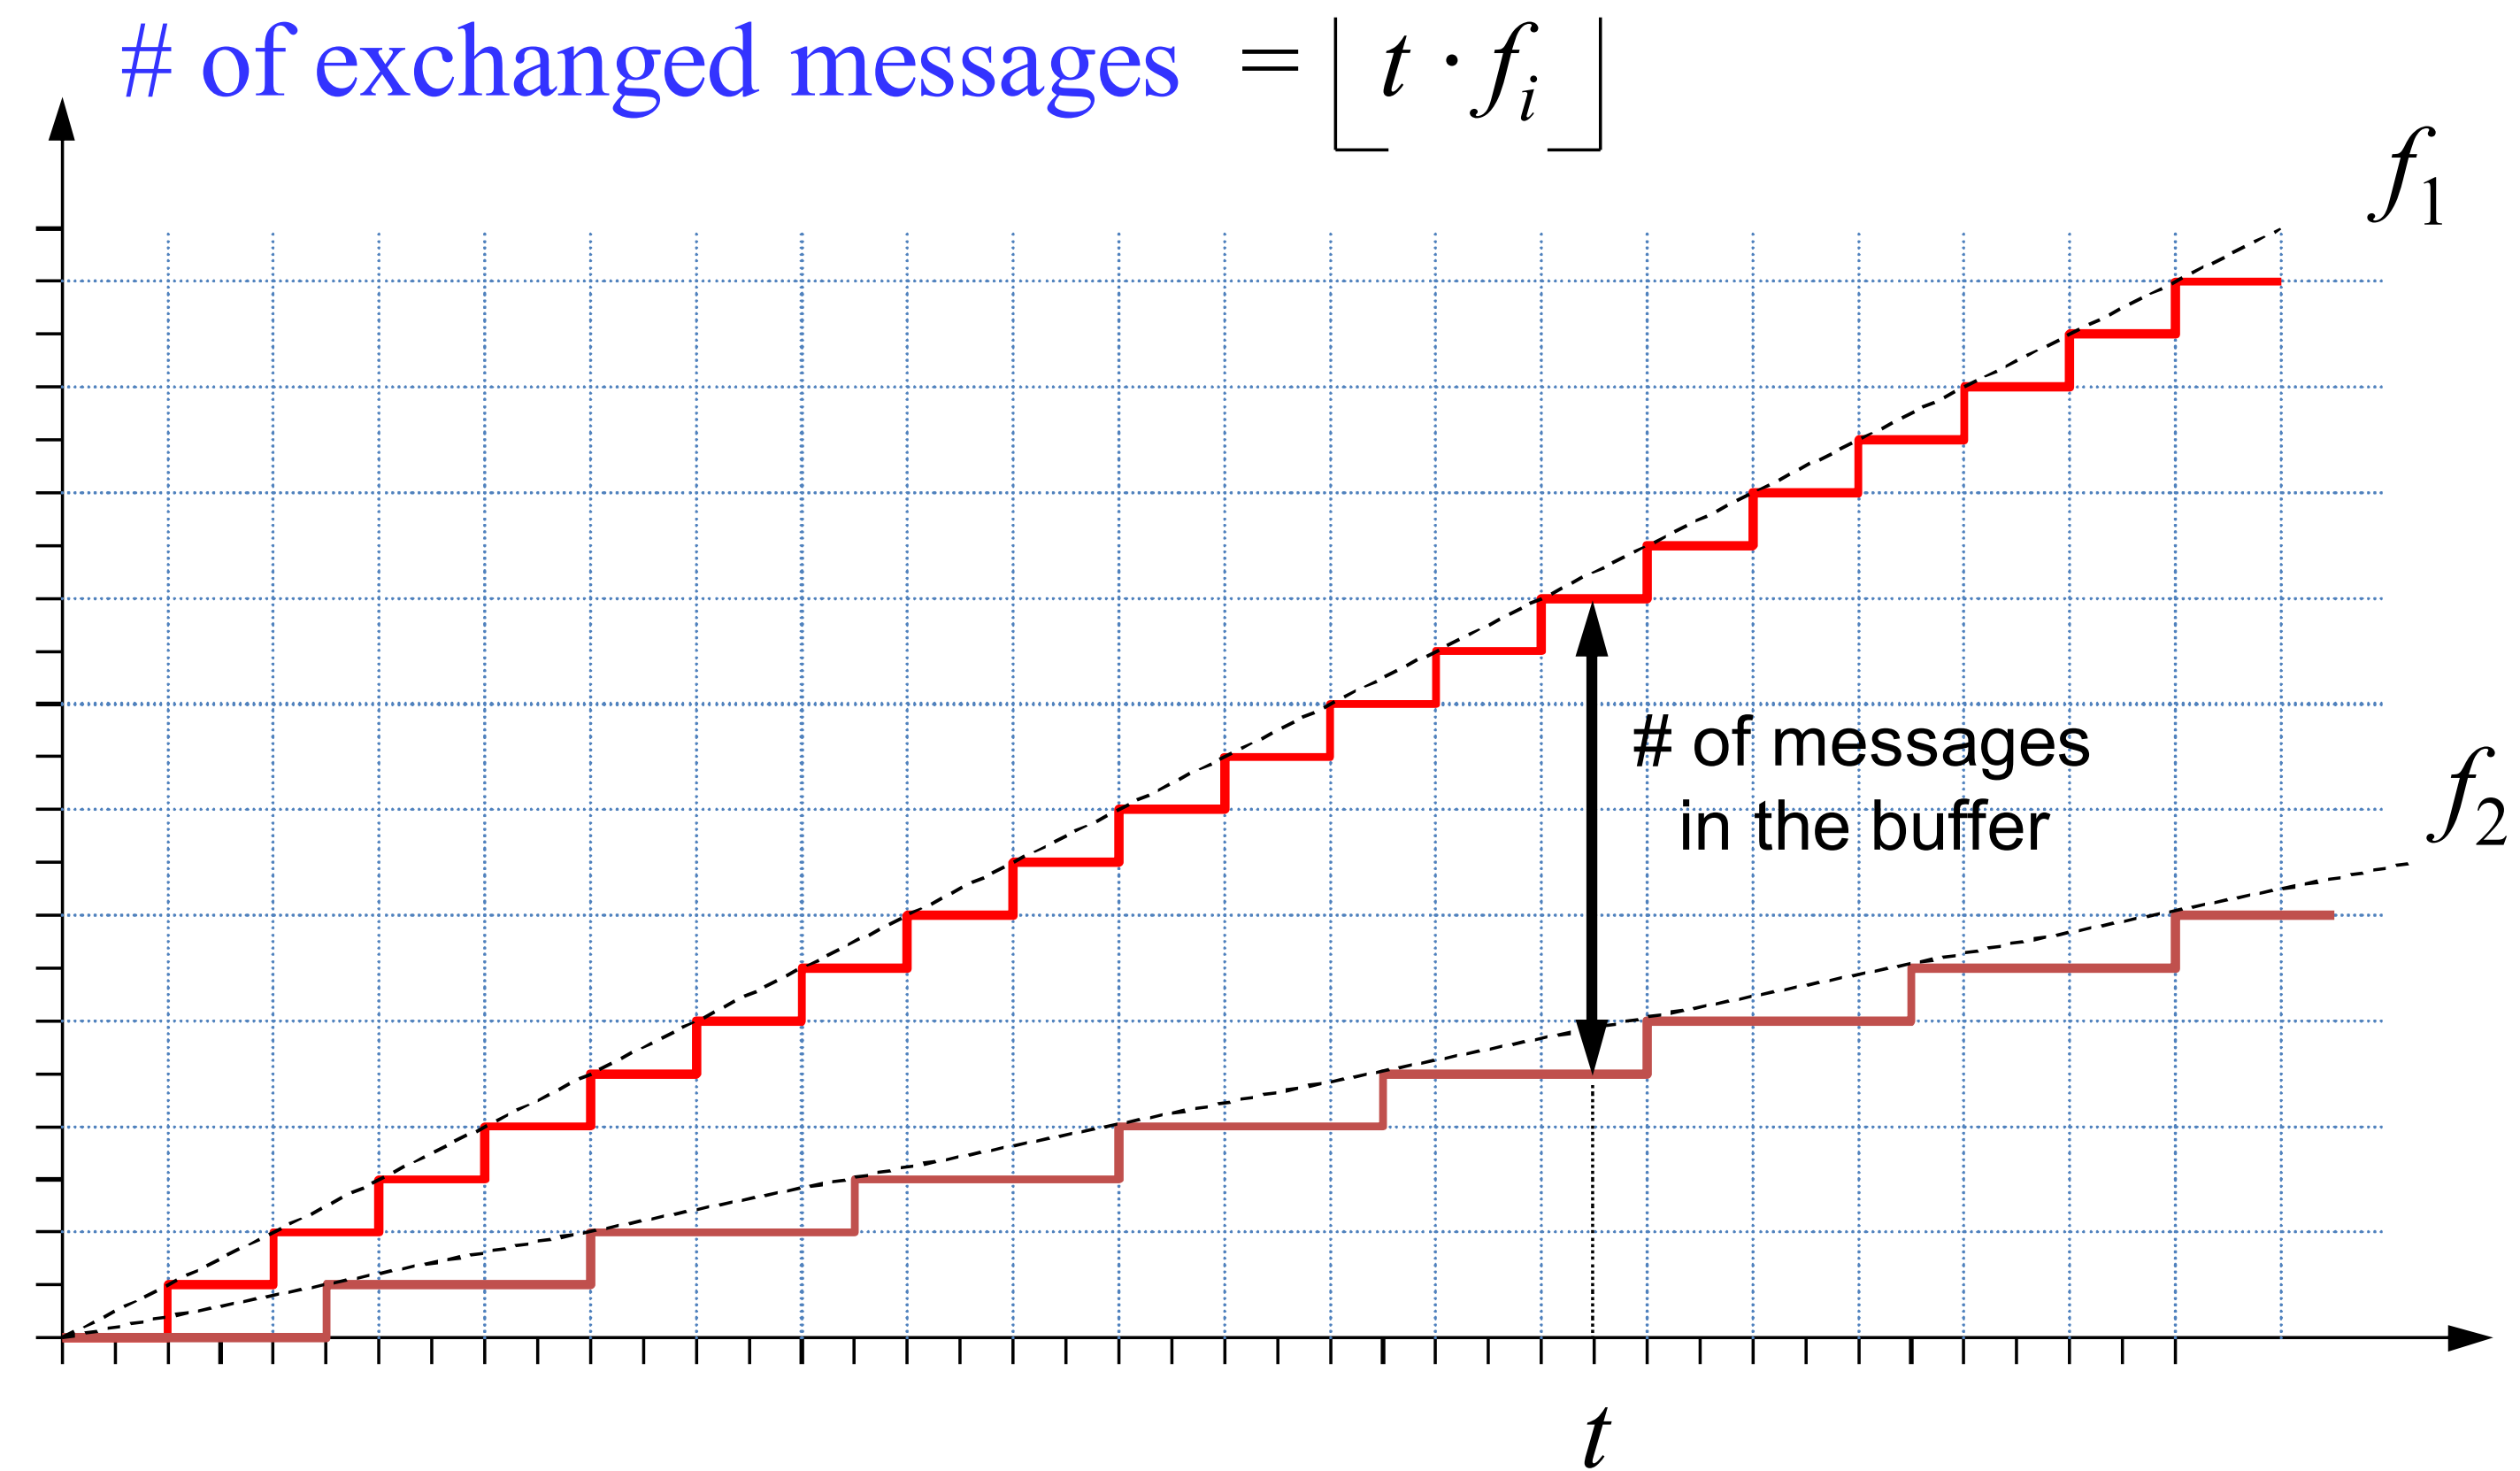
Thus, if  $T_1 \neq T_2$ , after a certain time, both tasks will proceed at the rate of the slowest task.



To keep its own rate, tasks using synchronous ports should have the same period.

How long two tasks with different periods can run at their own rate before they synchronize on a full (or empty) buffer?

# Exchanged messages



# Buffer saturation

If  $T_1 < T_2$ , the buffer saturates when:  $\left\lfloor \frac{t}{T_1} \right\rfloor - \left\lfloor \frac{t}{T_2} \right\rfloor > N$

Hence, the tasks proceed at their proper rate while:

$$\left\lfloor \frac{t}{T_1} \right\rfloor - \left\lfloor \frac{t}{T_2} \right\rfloor < \frac{t}{T_1} - \frac{t}{T_2} + 1 < N$$

That is, while:

$$t < (N - 1) \frac{T_1 T_2}{T_2 - T_1}$$

# STICK Ports

They are ports with a state message semantics:


- the **most recent message** is always available for reading;
- a new message **overrides** the previous one;
- a message is **not consumed** by the receiver.

## NOTE:


- Since only the most recent message is of interest, there is no need to maintain a queue of past messages.
- A task never blocks for a full or empty buffer.

# STICK Ports

**Example**  
 $T_1 < T_2$



**Example**  
 $T_1 > T_2$



# Blocking on STICK ports

- Although a task cannot blocks for a full or empty buffer, it can block for **mutual exclusion**.


Indeed, a semaphore is needed to protect the internal buffer from simultaneous accesses.

- Long messages may cause long blocking delays on such a semaphore.
- Long waiting times due to long messages can be avoided through a **buffer replication** mechanism.




# Dual buffering

Dual buffering is often used to transfer large data, (images) from the input peripheral device to a task:



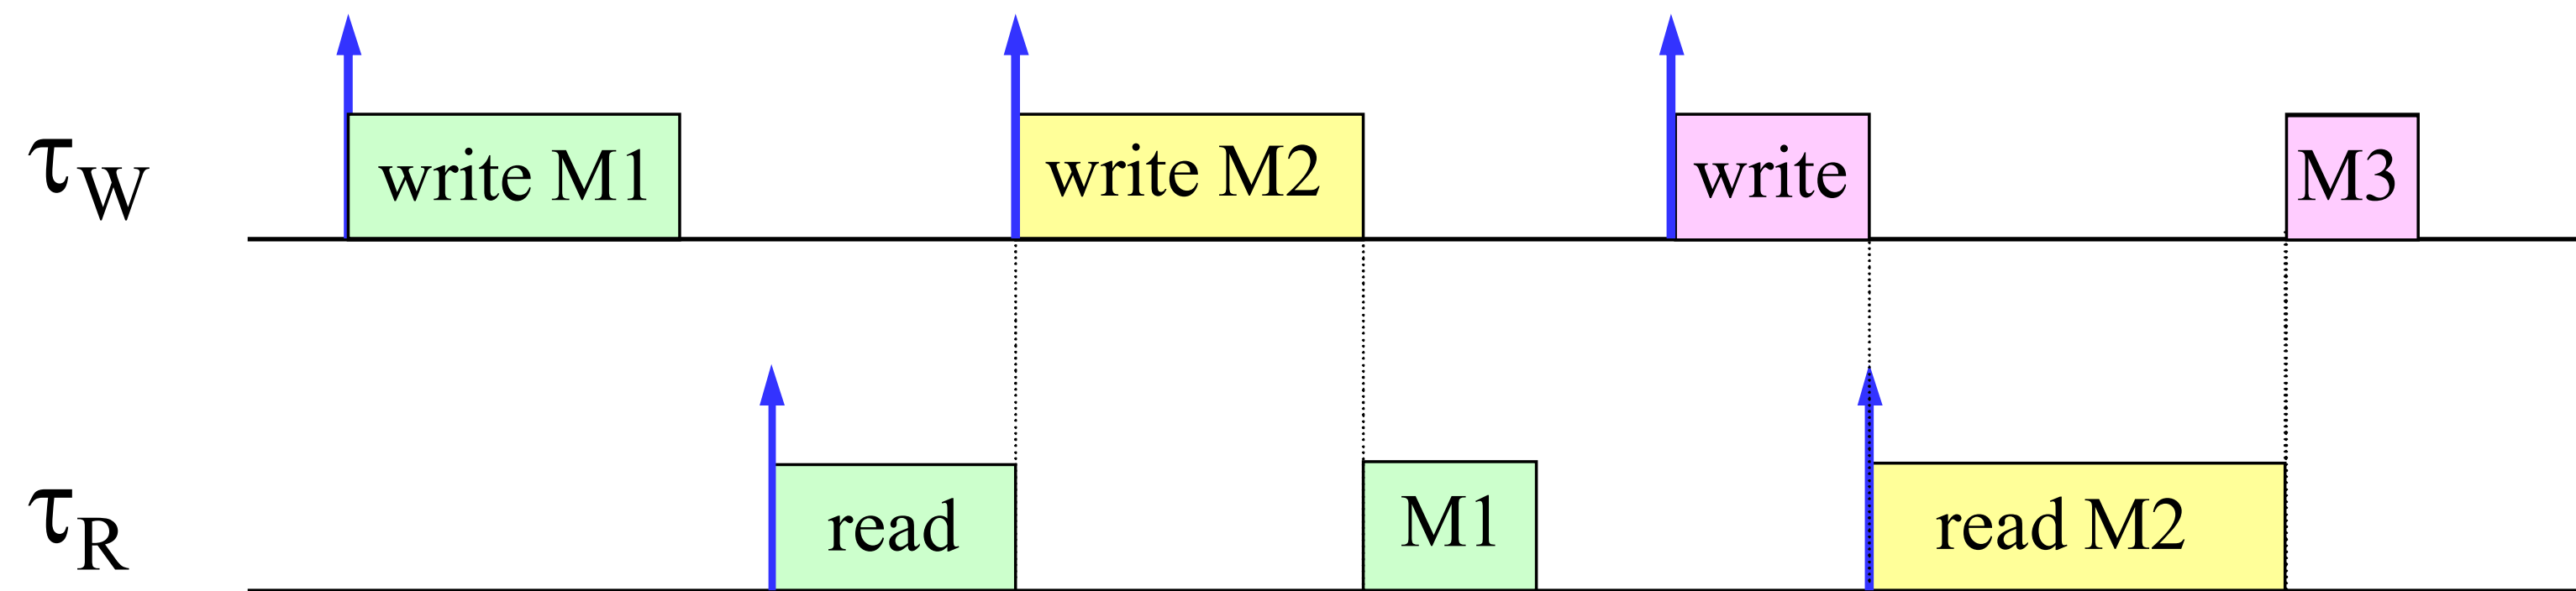
# Dual buffering

If a writer task  $\tau_w$  produces a new message while  $\tau_r$  is reading, the new message is written in a new buffer:



# Dual buffering

A new message becomes available to the next reader only when completely written:



# **Cyclic Asynchronous Buffers (CABs)**


# Cyclic Asynchronous Buffers

A **CAB** generalizes dual buffering for  $N$  readers:



- It is a mechanism for exchanging messages among periodic tasks with **different rates**.
- Memory conflicts are avoided by replicating the internal buffers.
- It uses a **state-message semantics**:
  - At any time, the most recent message is always available for reading;
  - A new message overrides the previous one;
  - Messages are not consumed by reading.

# Accessing a CAB

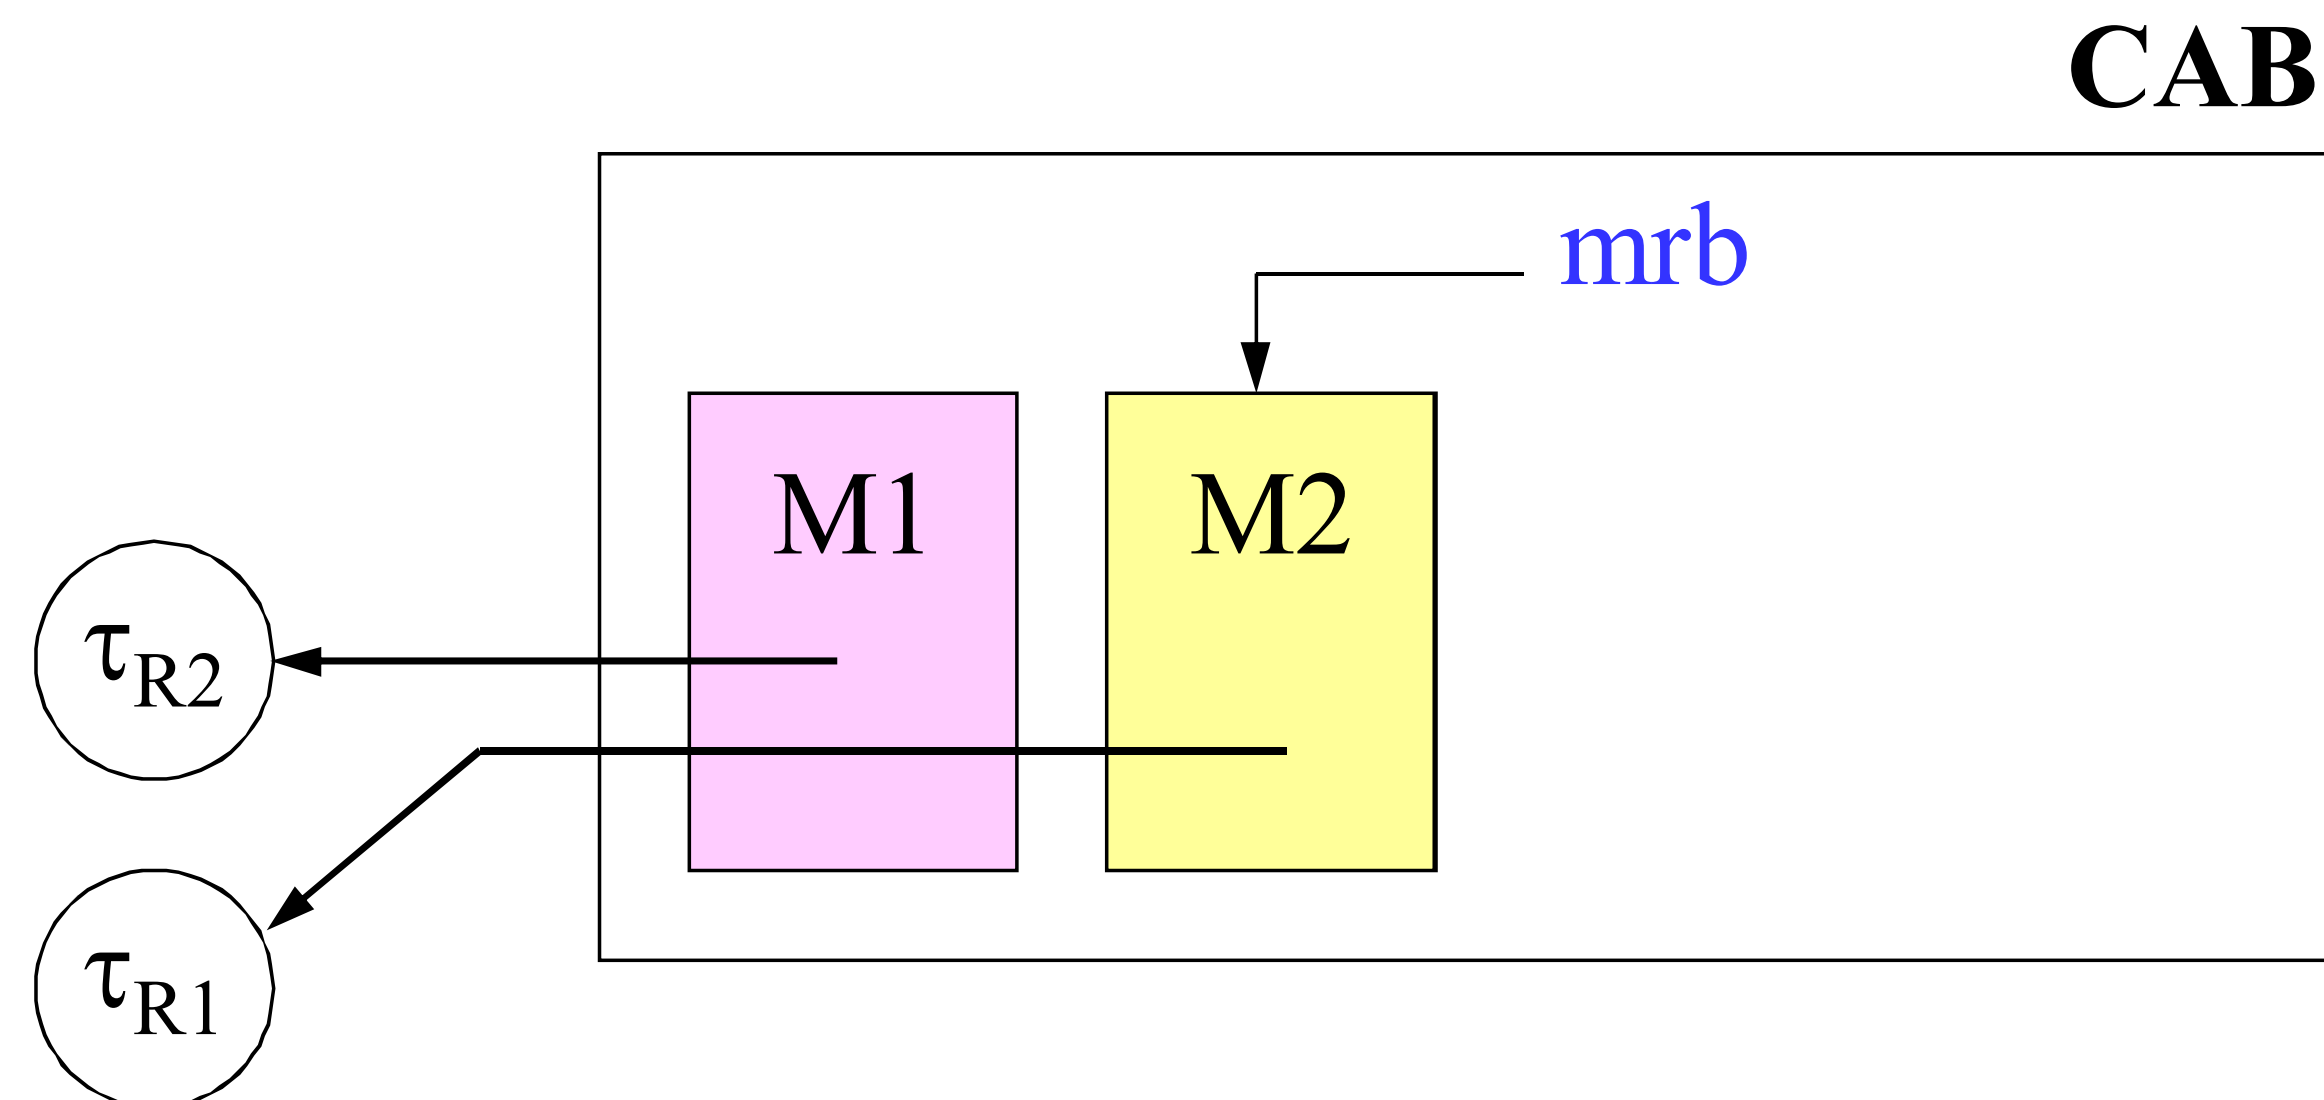

- A message is accessed through a **memory pointer**.
- At any time, the **most recent buffer** can be accessed through a pointer (**mrp**) available to any reader.
- Hence, a reader is not forced to copy the message in its memory space.
- More tasks can **concurrently** read the same message.



# CAB example

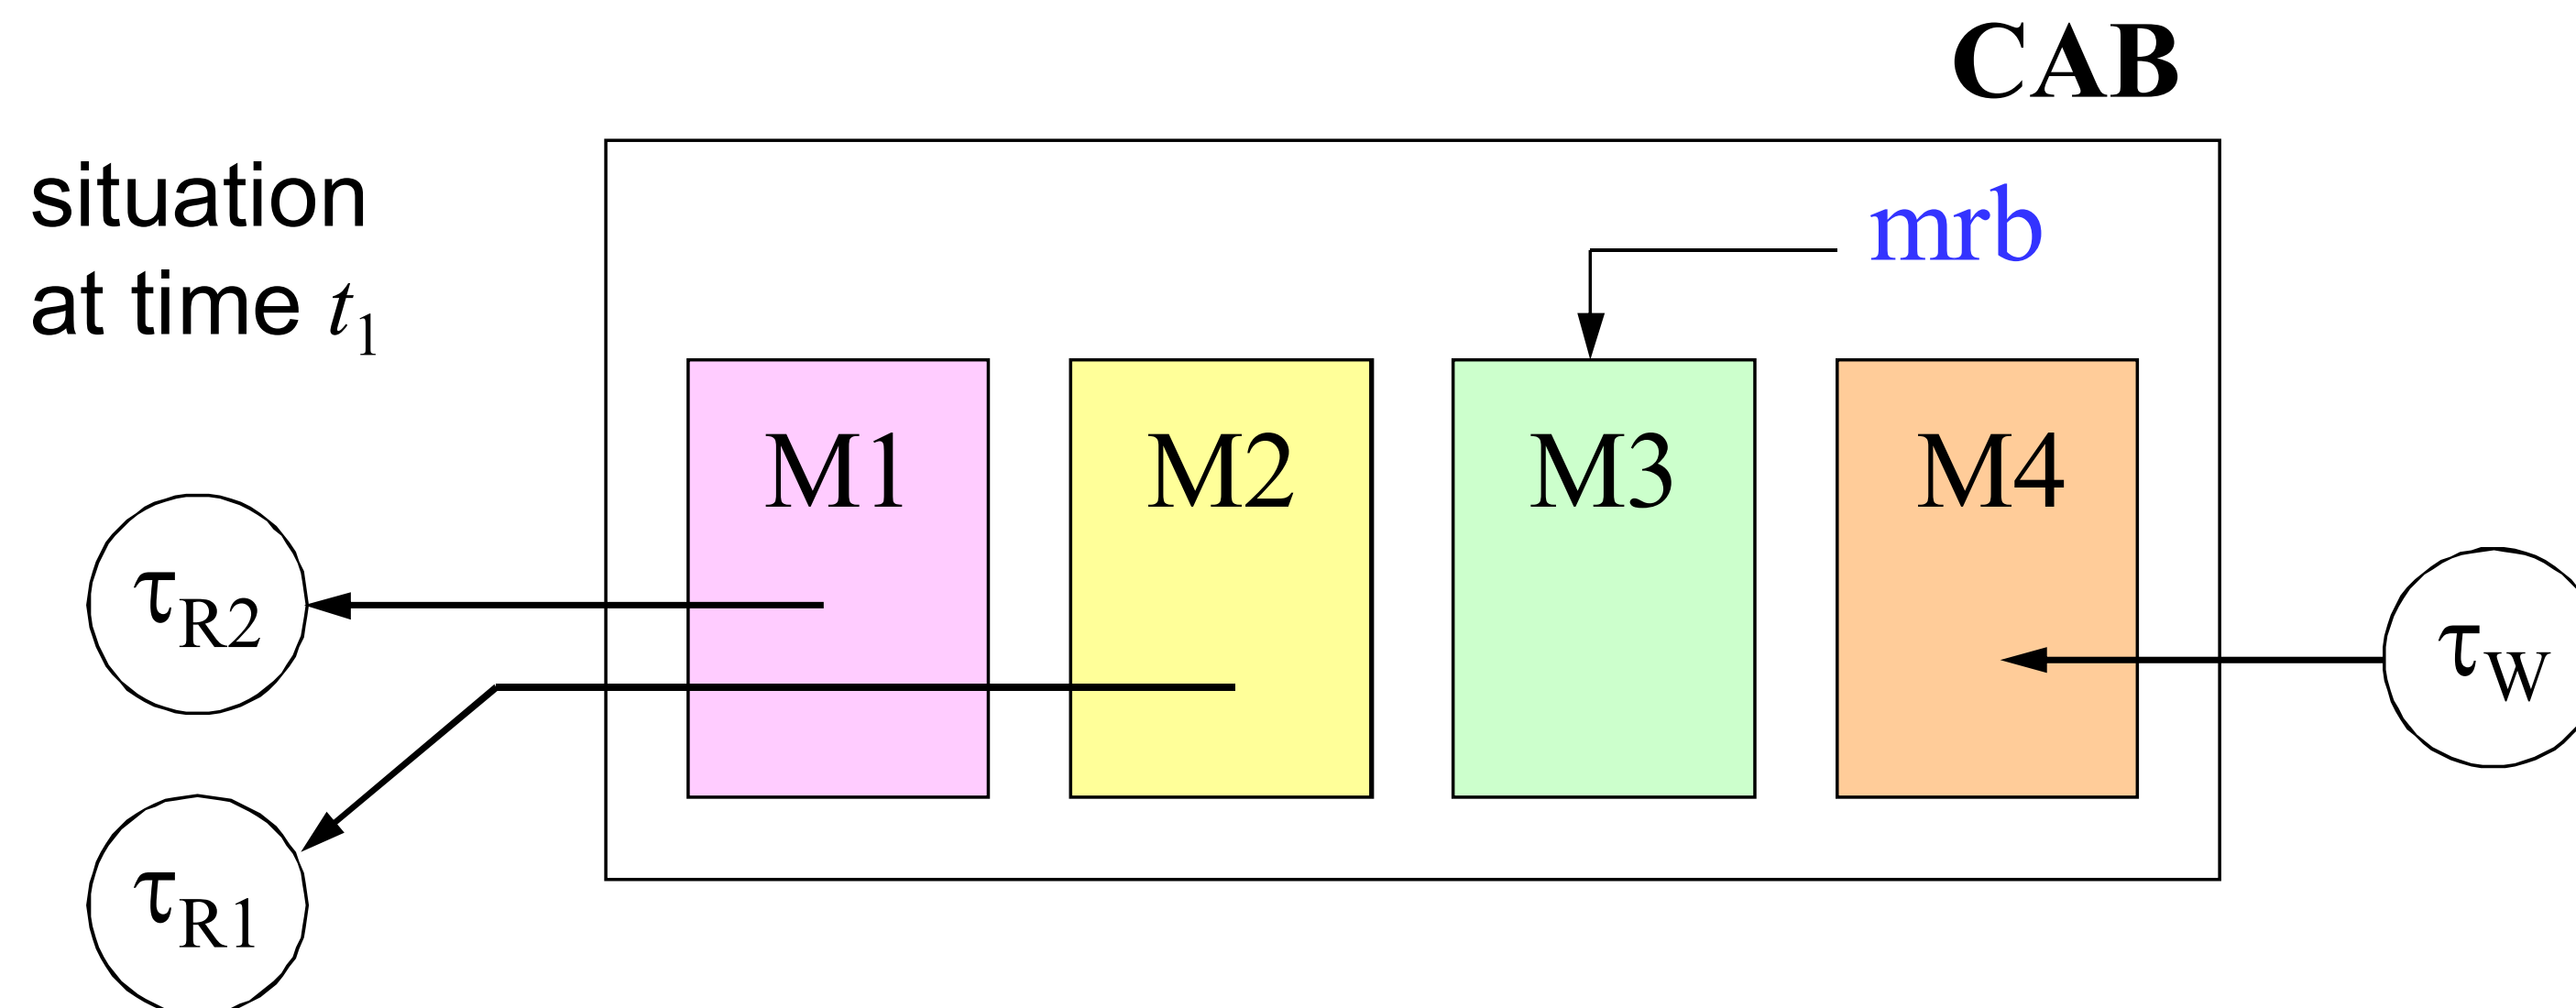



# CAB example

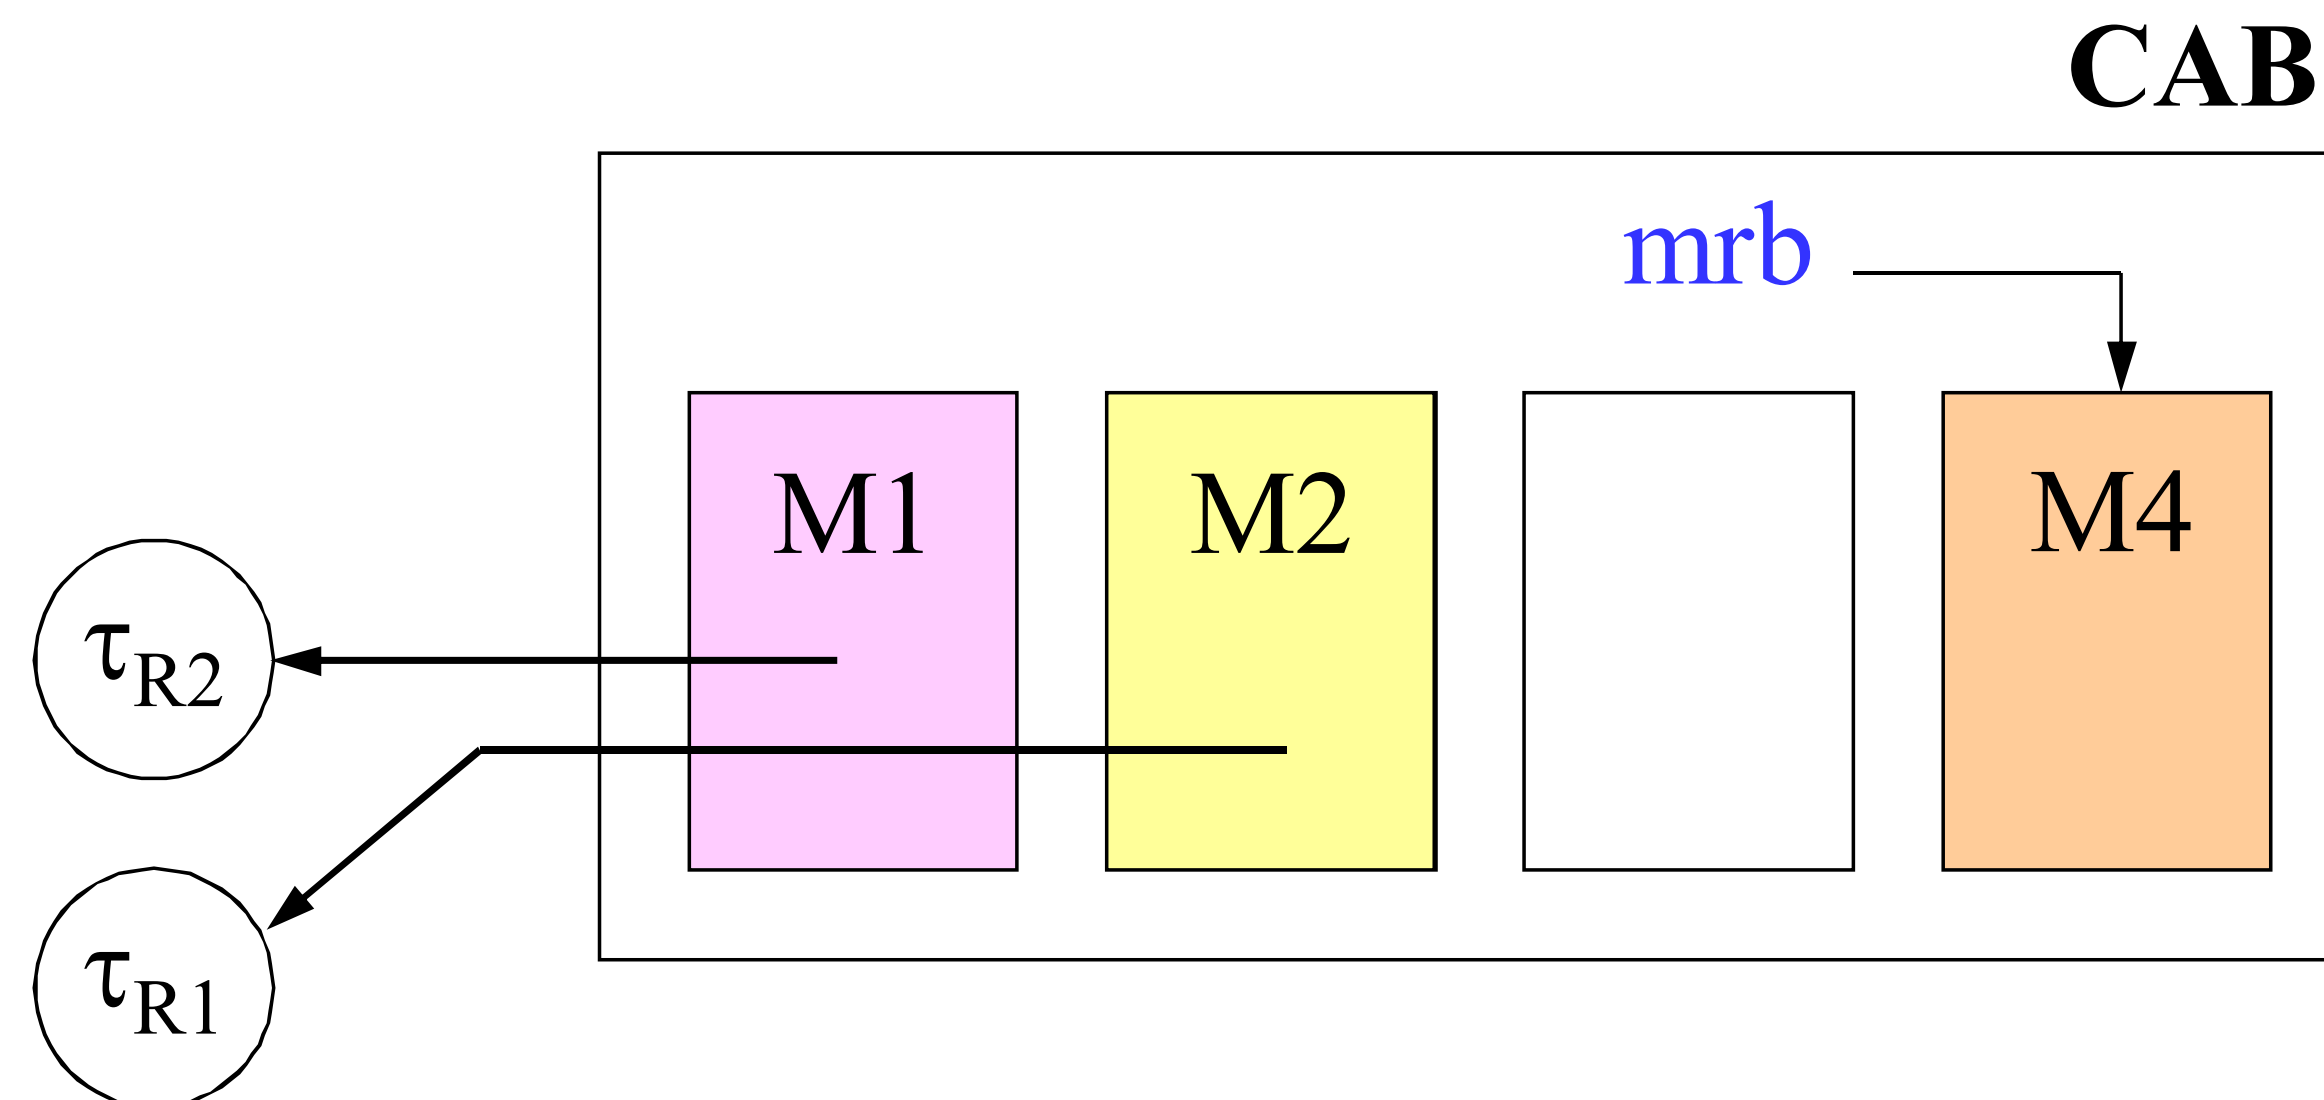
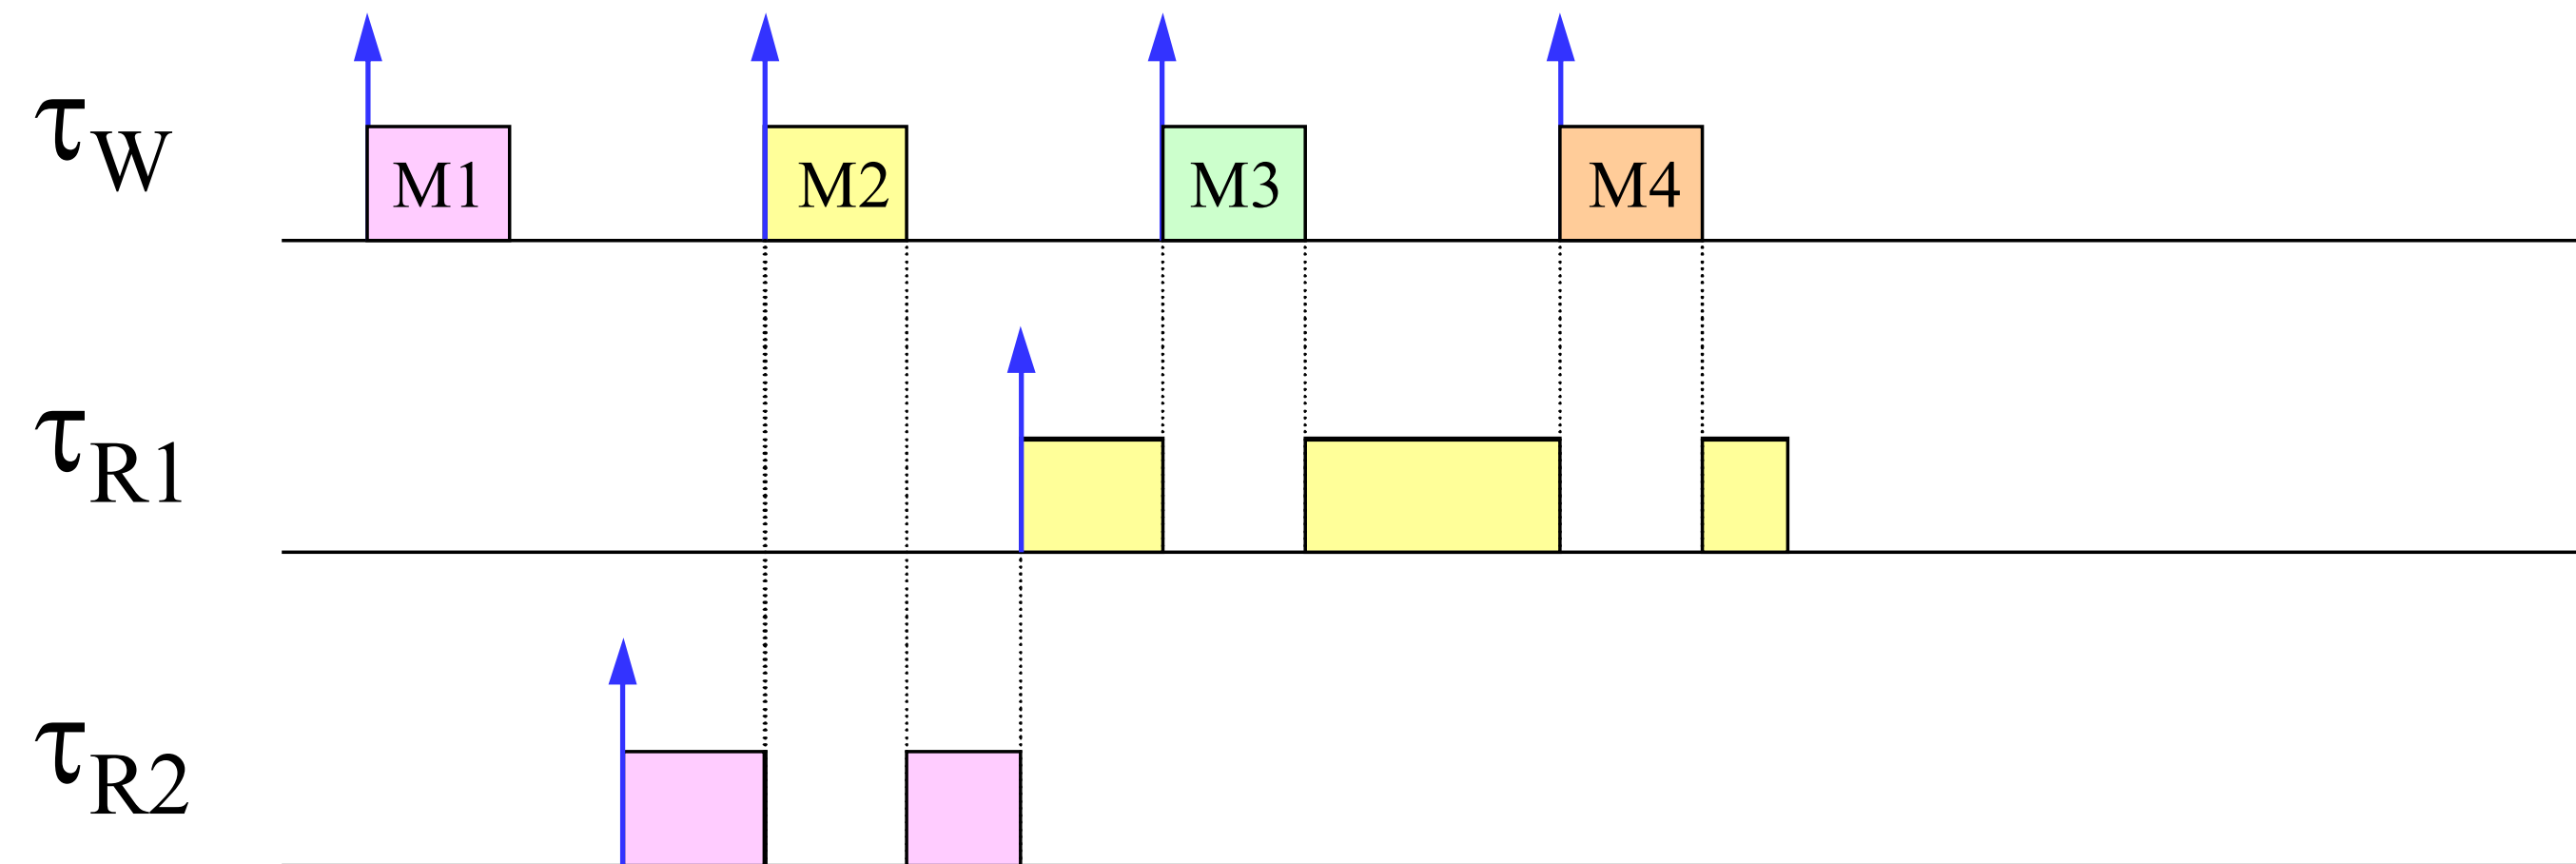





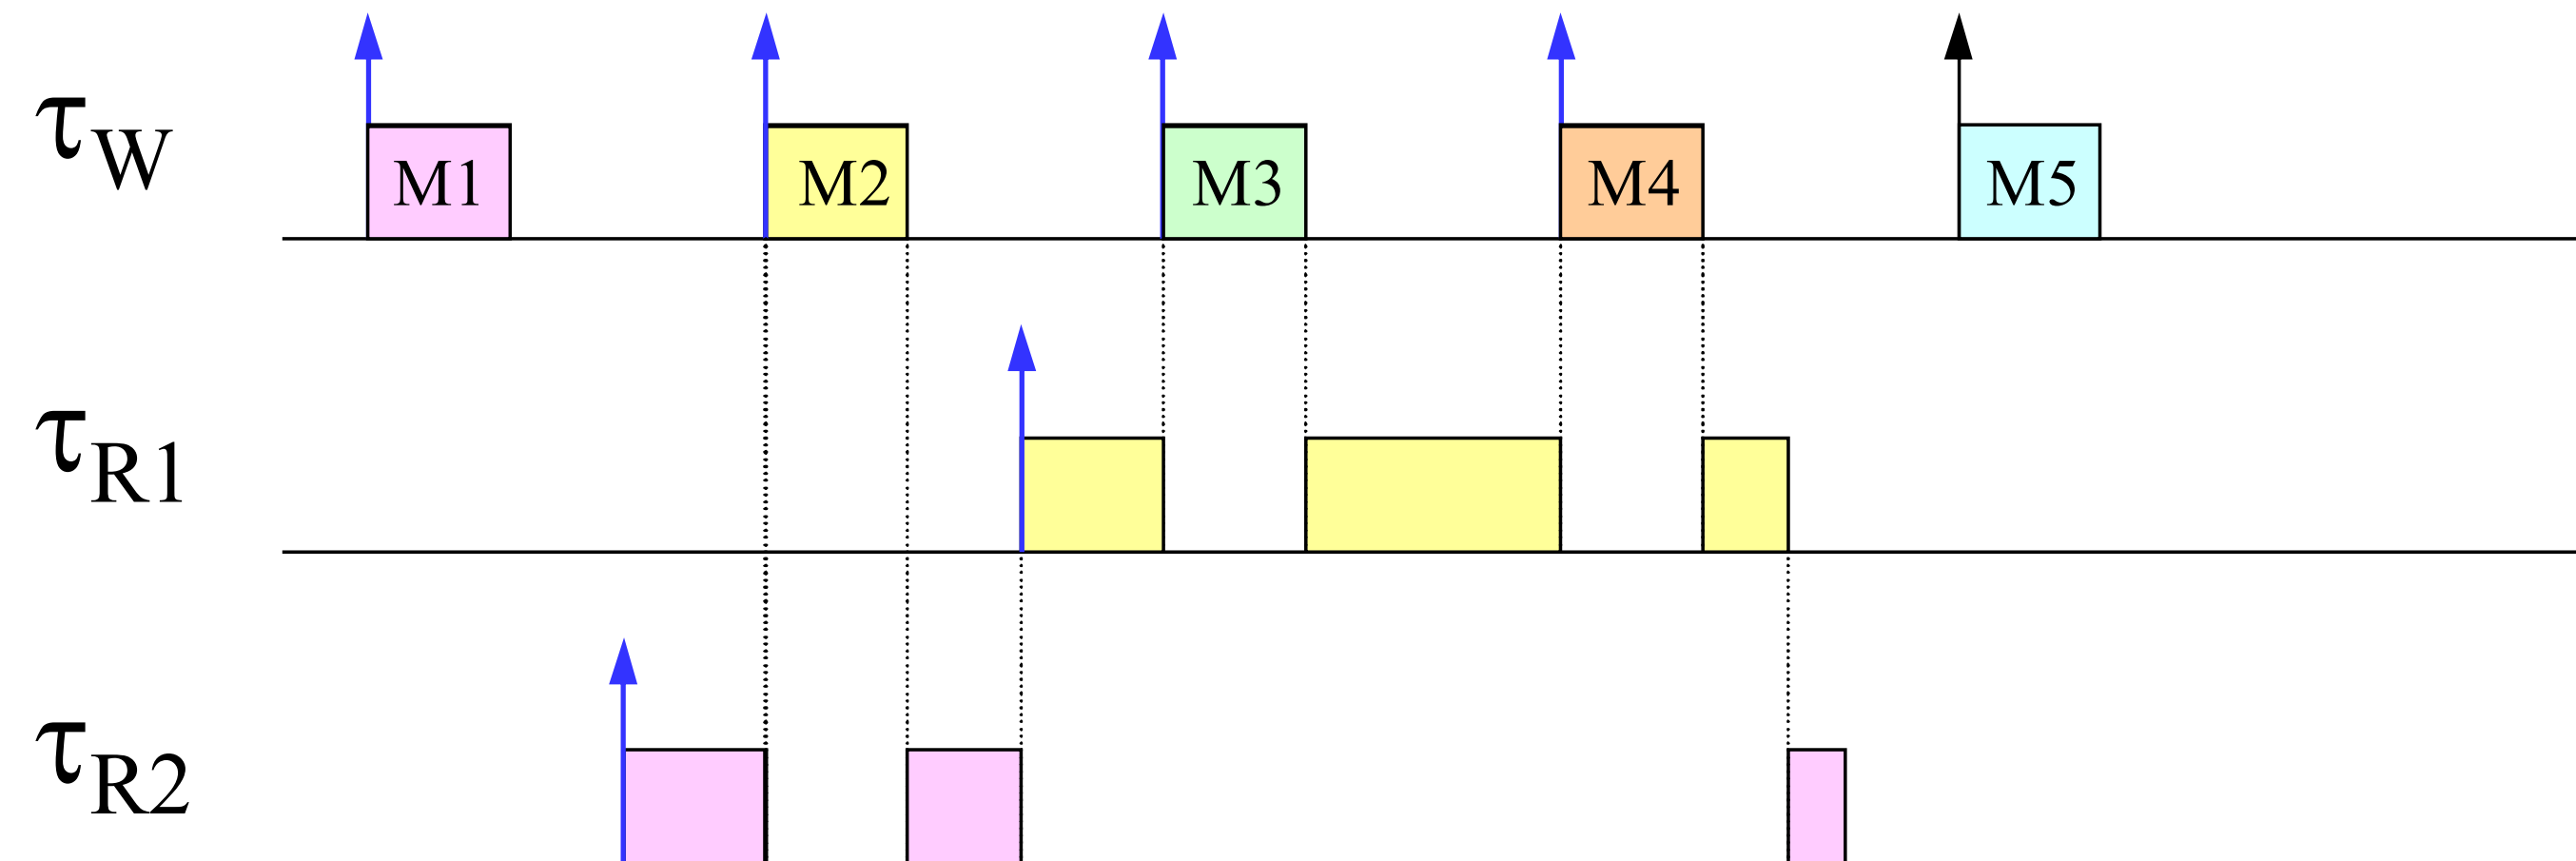
# CAB example




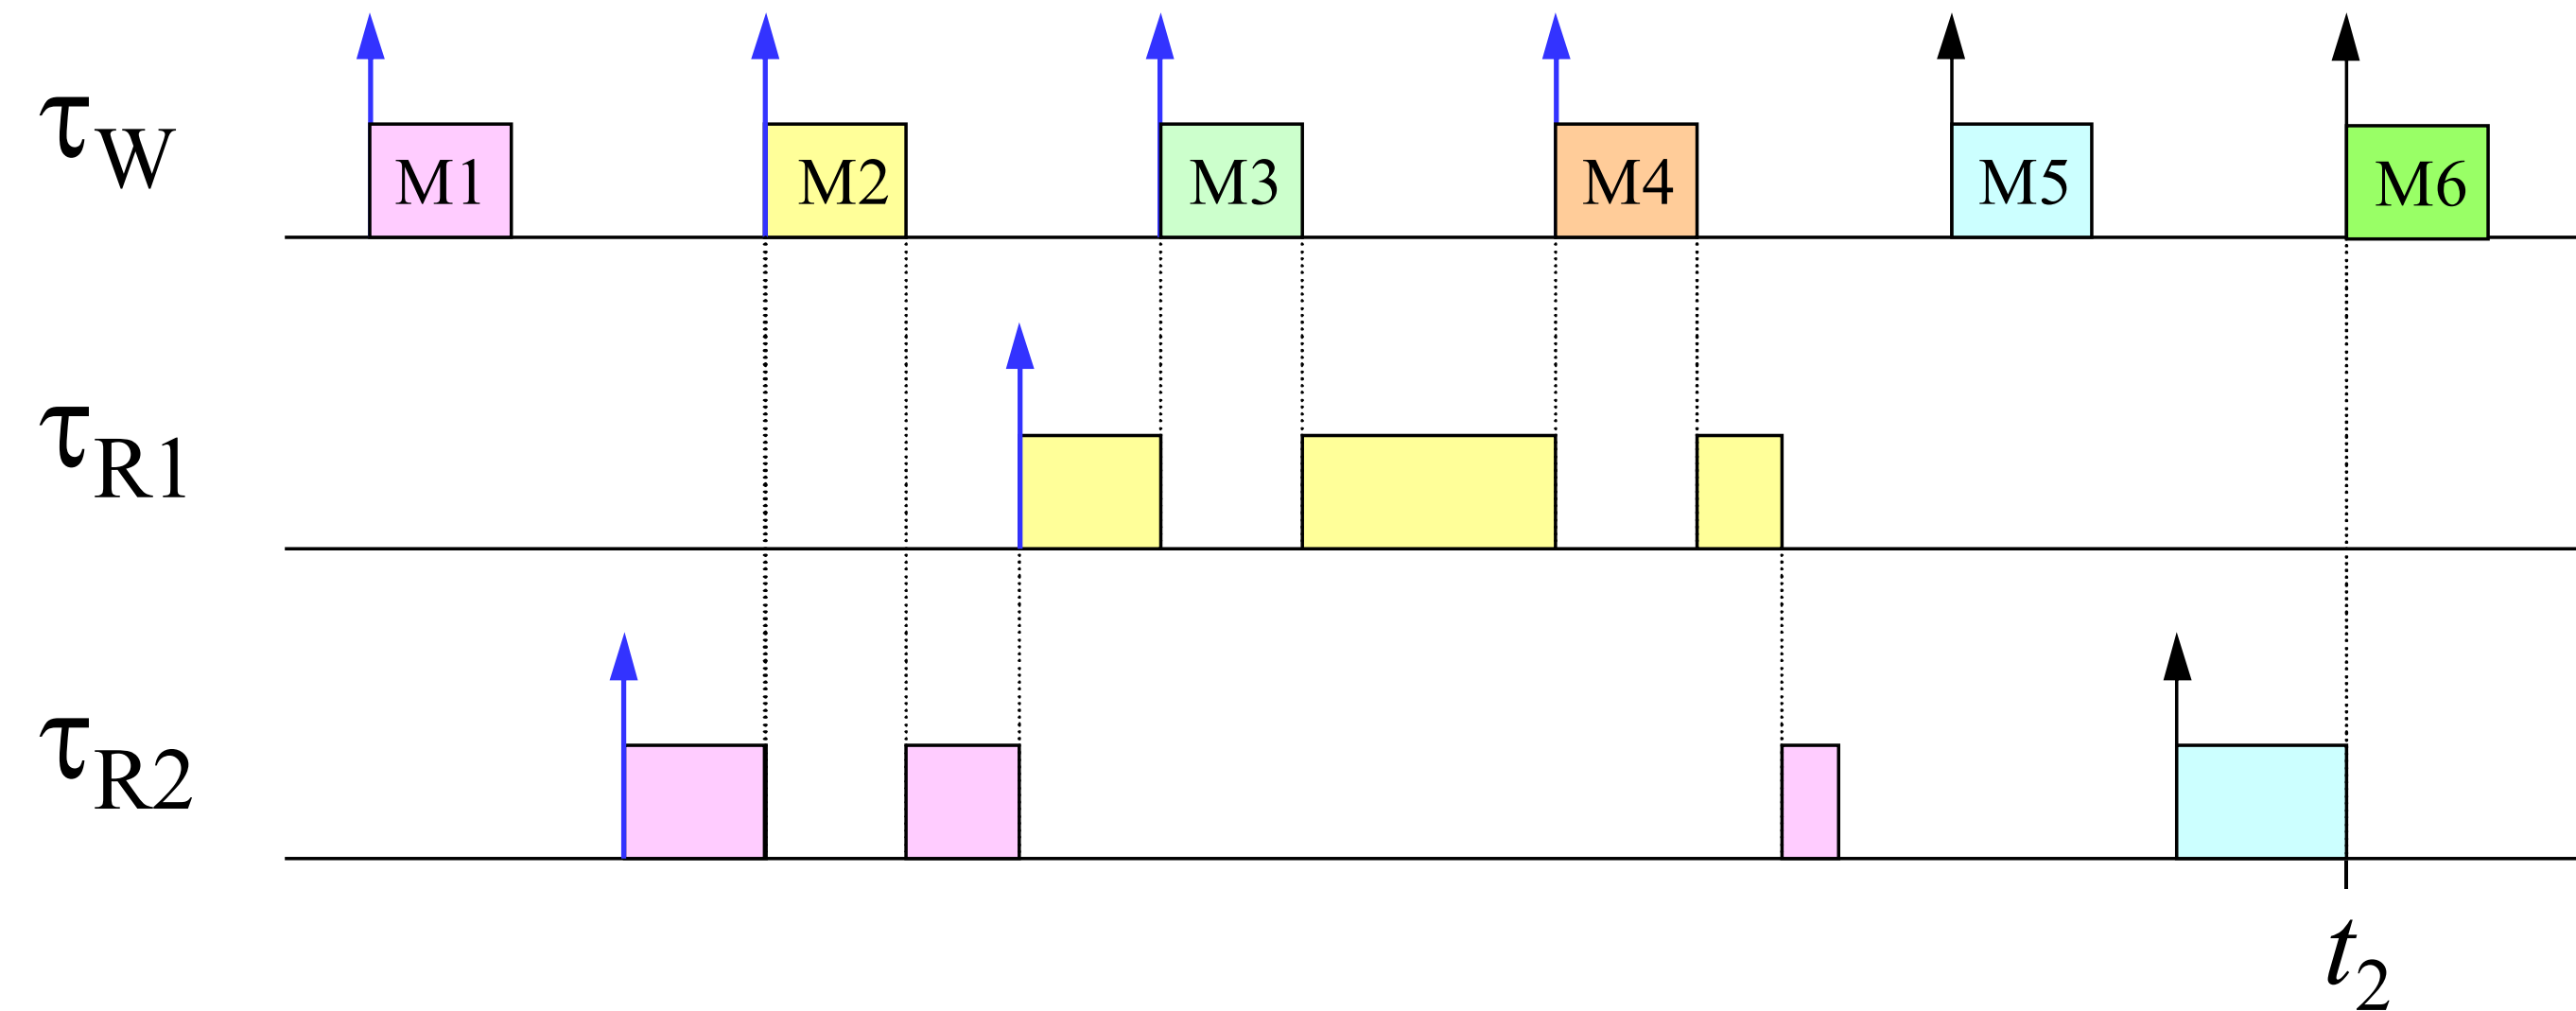
# CAB example



# CAB example



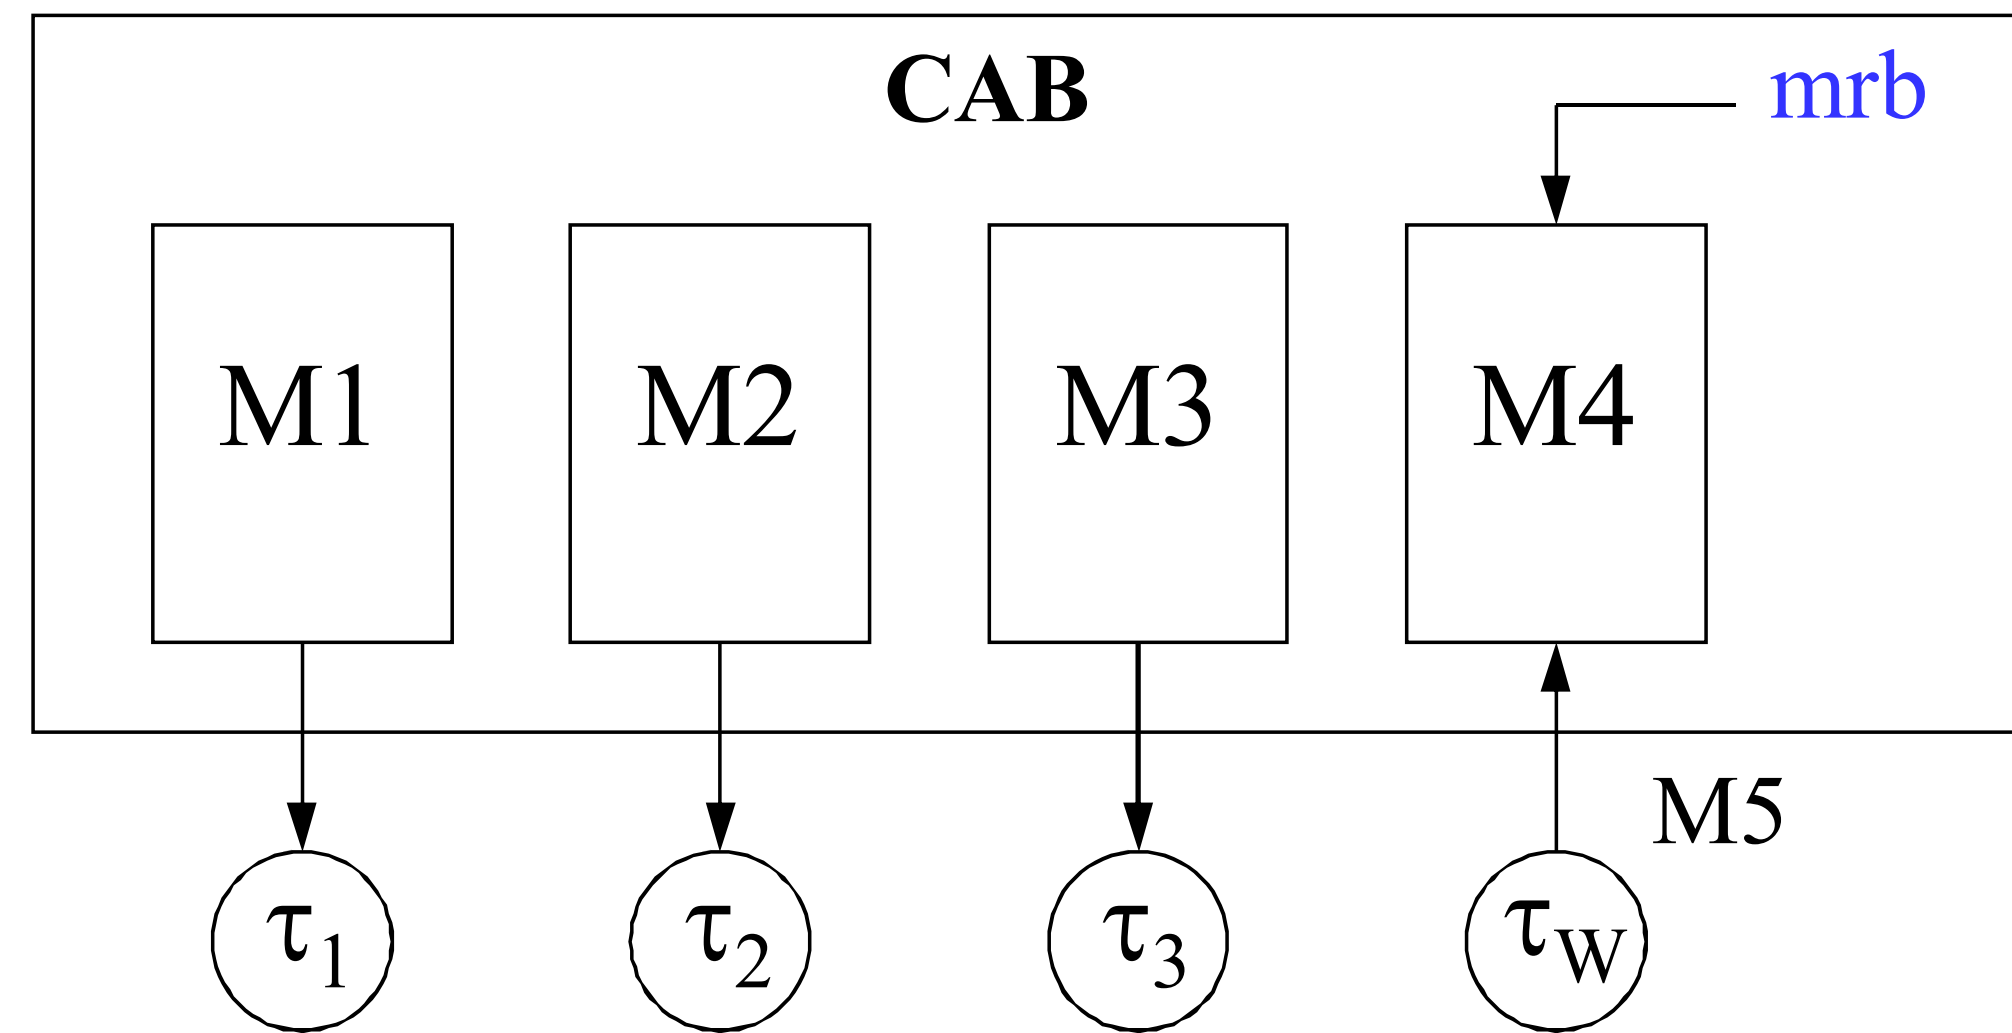
# CAB example



# Dimensioning a CAB

- If a CAB is used by  $N$  tasks, to avoid blocking, it must have at least  $N+1$  buffers.
- The  $(N+1)^{\text{th}}$  buffer is needed for keeping the most recent message in the case all the other buffers are used.

# Inconsistency with N buffers



- Assume that all the buffers are used and  $\tau_w$  overwrites the most recent message (M4) with M5.
- If (while  $\tau_w$  is writing)  $\tau_1$  finishes and requests a new message, it finds the CAB inconsistent.

# Writing Protocol

To write a message in a CAB a task must:

1. ask the CAB for a pointer to a free buffer;
2. copy the message into the buffer using the pointer;
3. release the pointer to the CAB to make the message accessible to the next reader.

# Reading Protocol

To read a message from a CAB a task must

1. get the pointer to the most recent message in the CAB;
2. process the message through the pointer;
3. release the pointer, to allow the CAB to recycle the buffer if it is not used.